

Apache Kafka

from zero
to semi hero



Fincons Tech Hub
Java Spring Time



Daniele Gallitelli
Software Eng. – Senior Analyst



Francesco Lacriola
Software Eng. – Senior Analyst



Michele Caccia
Senior Software Eng. – Senior Consultant

What we'll do today

- Introduction to Apache Kafka
- Build a real world, multi-service event-driven system
 - Simulates a fleet of trains sending real-time location data
 - Processes data to generate insights
 - Visualizes everything on a live dashboard
 - Uses multiple programming languages (Java & Python)
- Q/A



What is Apache Kafka

Distributed streaming platform for building real-time data pipelines



- Originally developed by LinkedIn, now open-source (Apache)
- **High-throughput message broker**, allowing systems to publish, store, and consume **streams of records** in a **fault-tolerant and scalable** way
- **Durable storage** — Data is persisted on brokers disk, allowing replay and recovery.

Key Features


- ✓ High throughput & low latency
- ✓ Fault-tolerant & scalable
- ✓ Persistent storage
- ✓ Real-time processing



What is Apache Kafka



What is Apache Kafka



Where is Kafka used

Used by 2000+ firms, 80% of Fortune 100

- **SBB:** Yes — we use Kafka too! (Master Data propagation, connected train telemetry and real-time services)
- **Netflix:** Real-time recommendations & monitoring
- **Uber:** Real-time pricing, trip tracking
- **LinkedIn:** Activity streams, operational metrics
- **PayPal:** Risk detection, fraud prevention
- **Banking:** Transaction processing, fraud detection
- **E-commerce:** Inventory management, order processing



Topics: Your Data Channels

Topic = A particular stream of data published in your Kafka Cluster

- Like a table in a database or a folder in a filesystem
- Topics are **multi-subscriber** (many applications can read the same topic)
- A topic is identified by its name
- It is impossible to query topics. Use Kafka Producers to send data and Kafka Consumers to read the data

Kafka Broker

logs

train-locations

train-speed
average


Examples in our workshop:

- train-locations - raw position data from trains
- train-speed-averages - processed speed analytics



Partitions and offsets

- Each topic is divided into **partitions**
- Partitions allow parallel processing
- Messages in a partition can be **ordered**
- Each message has an **offset** (sequential ID)
- Kafka topics are **immutable**
- Our train-locations topic has **3 partitions**




```
Topic: train-locations
└─ Partition 0: [msg0, msg3, msg6, ...]
└─ Partition 1: [msg1, msg4, msg7, ...]
└─ Partition 2: [msg2, msg5, msg8, ...]
```



Producers: Publishing Data

Applications that **publish** (write) records to topics

- Specify which partition to send to (and which Kafka broker has it)
- Can use **keys** for consistent routing
- In our workshop: **TrainLocationSimulator** (Java)



Example:

Train T-123 → partition based on train ID → always same partition

This ensures all messages from the same train stay in order!


Consumers: Reading Data

Applications that **subscribe** to topics and process records (pull model)

- Maintain an **offset** to track what they've read -> Consumers know which broker to read from
- Can rewind and replay messages

In our workshop:


- **SpeedAnalysisStream** (Kafka Streams)
- **DashboardWebApp** (Java)
- **MaintenanceAlerter** (Python)



Consumer Groups: Teamwork!


Multiple consumers working together as a **group**

- Each partition assigned to **only one** consumer in the group
- Enables **parallel processing** and **load balancing**
- If a consumer fails, its partitions are reassigned → **fault tolerance**



Benefits:

- ✓ Same group = share the workload
- ✓ Different groups = each gets all messages
- ✓ Zero message loss



notification service

train-location service


Schema Registry

Apache Kafka doesn't validate data streams

- What if a producer sends bad data?
- What if a field gets renamed?
- What if the data format changes from one day to another?



consumers breaks!



Apache Avro

Data Serialization Framework

efficient, compact, binary
format

Schema Evolution

backward & forward
compatibility

Readability

schema readable (JSON)
Data is not (binary)

Schema-based

JSON-defined schemas &
ensures data consistency



Kafka vs RabbitMQ/ActiveMQ



Kafka Advantages:

- Higher throughput (millions of msgs/sec)
- Persistent storage (replay capability)
- Horizontal scalability
- Built for streaming & big data

Traditional Messaging:

- Complex routing (exchanges, bindings)
- Lower latency for small messages
- Message priority queues
- Simpler request/reply patterns

When to use Kafka:

- High volume event streaming
- Log aggregation
- Real-time analytics
- Event sourcing



Our Workshop Architecture




Services:

1. **Producer:** TrainLocationSimulator (Java) - generates train positions
2. **Maintenance Alerter:** MaintenanceAlerter (Python) - slow train alerts
3. **Dashboard Consumer:** TrainConsumerDashboard (Java) - Demonstrates partition rebalancing
4. **Stream Processor:** SpeedAnalysisStream (Java, Kafka Streams) - calculates averages
5. **WebSocket Dashboard:** TrainWebsocket (Java, WebSocket) - real-time visualization


All connected through Kafka topics!



Our Workshop Architecture: topic train-locations



Our Workshop Architecture: topic train-speed-averages



Kafka is Language Agnostic



Any language can produce/consume from Kafka:

- ➊ Java (most common, best support)
- ➋ Python
- ➌ JavaScript/Node.js
- ➍ Go
- ➎ C#/.NET

And many more...

Workshop Demonstration:

- ✓ Java producers and consumers
- ✓ Python consumer (maintenance alerter)
- ✓ Both read from the **same topic** seamlessly



Kafka Streams



Library for building stream processing applications

- Processes data **as it arrives** (real-time)
- Supports transformations, aggregations, joins, windowing
- No separate cluster needed

Our Use Case:

1. Read from train-locations
2. Calculate average speed per train (10-second window)
3. Write to train-speed-averages



Kafka Monitoring Tools



Command Line Tools:

```
● ● ●  
kafka-topics --list  
kafka-topics --describe --topic train-locations  
kafka-consumer-groups --describe --group dashboard-group
```

Web Interface:

- **Kafka UI** (localhost:8099, shipped with our containers)
 - Monitor topics, partitions, consumer groups
 - View messages in real-time
 - Track consumer lag
- **RedPanda**
- ... and many more!



Common Patterns

Event Sourcing

- 📝 Store every change to application state as an immutable sequence of events.
- 🧩 Kafka acts as the **event log** — a single source of truth.
- ⌚ Enables **time-travel debugging**, replay, and rebuilding state at any point.

CQRS (Command Query Responsibility Segregation)

- 👉 Separate **commands (writes)** from **queries (reads)** for scalability and clarity.
- 💡 Kafka Streams can manage the **write-side events**, while **materialized views** serve reads.
- ⚖️ Perfect for systems with **uneven read/write loads**.



Common Patterns

Data Pipelines

- 🔗 Kafka connects **producers and consumers** in real-time data flow.
- ⚙️ Simplifies ETL: **ingest → process → distribute**, no batch delays.
- 📊 Powers **streaming analytics** and long-term data storage.

Microservices Communication

- ✉️ Kafka decouples services with **asynchronous, event-driven** messaging.
- 🧠 Reduces reliance on REST calls, improving **resilience** and **fault tolerance**.
- 🌱 Promotes **loose coupling** and **scalable** architecture.



Production Best Practices



Design:

- 📦 Use shared libraries for data models (avoid duplication)
- 🗄️ Schema Registry for data contracts (Avro, JSON Schema)
 - 📌 Check avro services in the repository for more!
- 🔄 Version your message formats
- 🎯 Use message keys for partitioning logic (same key = same partition)
- ⚡ Configure appropriate retention policies

Execution:

- 📊 Monitor consumer lag
- ⚖️ Right-size partition count
- 🔒 Enable authentication & encryption (production)
- 💾 Configure replication factor
- 🧪 Test rebalancing scenarios



Hands-On Time! 🛠

What You'll Do:

1. Set up Kafka infrastructure (Docker)
2. Start the train location producer
3. Run the stream processor
4. Launch the WebSocket dashboard
5. Add Python maintenance alerter
6. Test partition rebalancing with scaling

Prerequisites check:

Docker & Docker Compose

If you want to build and run it yourself:

Java JDK 17+ • Maven • Python 3.8+



Partition Rebalancing in Action



Live Scaling Demo:

```
# Start with 1 consumer  
docker-compose up -d  
  
# Scale to 3 consumers  
docker-compose up -d --scale dashboard-consumer=3  
  
# Scale to 2 consumers  
docker-compose up -d --scale dashboard-consumer=2  
  
# Back to 1  
docker-compose up -d --scale dashboard-consumer=1
```

Watch: Partition reassignment • Rebalancing events • Zero message loss



Avoiding Common Mistakes !

Top Issues:

1. **Connection refused:** Kafka not running → docker ps
2. **Deserialization errors:** Mismatched serializers
3. **Consumer lag:** Too few consumers for partition count
4. **Wrong consumer group:** Competing consumers
5. **Offset reset issues:** Check auto.offset.reset config

Debug Commands:

```
● ● ●  
docker-compose logs kafka  
kafka-consumer-groups --describe --group <group-id>
```



Continue Your Kafka Journey



Deep Dive Topics:

- Advanced Kafka Streams (joins, state stores)
- Security (SSL, SASL, ACLs)
- Multi-datacenter replication
- Performance tuning
- Schema evolution strategies

Learning Resources:

- [Confluent Documentation](#)
- "Kafka: The Definitive Guide" (O'Reilly)
- Confluent Developer courses
- Apache Kafka official docs



Q/A



Thank You!

Workshop repository: [GitHub](#)

Call to Action:

Star the repository

Try building your own use case

Share your experience

Happy Streaming!

